Provenance-Aware Agentic Knowledge Graph Systems for Reliable Multi-Source Grounded AI

Primary Topic Area: Knowledge Grounding **Secondary Topic Area:** Agentic AI

Large Language Models (LLMs) increasingly rely on retrieval mechanisms to incorporate external knowledge during inference. Most current retrieval-augmented generation (RAG) systems rely on vector similarity search over unstructured documents Lewis et al. [2020]. While effective for relevance, these systems provide limited support for structured reasoning, provenance tracking, and verifiable evidence chains. As agentic AI systems become capable of autonomous planning and decision-making Park et al. [2023], this lack of grounded traceability introduces risks including hallucinated outputs Ji et al. [2023], unverifiable reasoning, and brittle performance when operating across heterogeneous knowledge sources.

This project proposes **Provenance-Aware Agentic Knowledge Graph Systems (PA-AKG)**, a hybrid architecture that integrates knowledge graphs Hogan et al. [2021], multi-source retrieval, and agent orchestration to support reliable grounded reasoning. The core idea is to combine structured knowledge representations with generative models so that agent workflows can reason over entities, relationships, and constraints while maintaining explicit links to supporting evidence. In this framework, agents dynamically select between vector-based retrieval and structured knowledge graph queries, enabling hybrid reasoning workflows that balance semantic retrieval with symbolic structure.

The proposed system enables **multi-source knowledge grounding** across knowledge graphs, relational databases, scientific literature, and domain-specific document collections. Agents generate reasoning chains that explicitly reference structured entities and supporting evidence. A key component is an automated **provenance generation and validation layer** that produces traceable reasoning paths linking model outputs to underlying knowledge sources, enabling both automated and human verification of agent outputs.

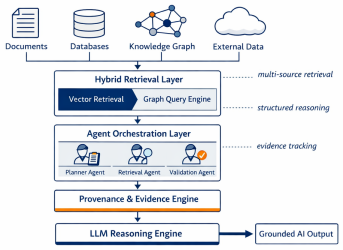

Figure 1: Hybrid Knowledge-Grounded Agentic AI Architecture. The PA-AKG architecture integrates heterogeneous knowledge sources with agentic reasoning workflows. A hybrid retrieval layer combines vector search and knowledge graph queries. An agent orchestration layer coordinates planning, retrieval, and validation agents. A provenance engine records evidence links enabling traceable verification of generated outputs.

This research investigates: (1) how heterogeneous knowledge sources can be integrated into a unified grounding framework for generative systems; (2) how agent workflows should dynamically select between structured and unstructured retrieval pathways; (3) what mechanisms enable scalable provenance generation and verification; and (4) how structured grounding affects reasoning accuracy, robustness, and efficiency relative to conventional RAG systems.

We will design a **hybrid KG–RAG architecture**, implement agent orchestration workflows combining structured reasoning with generative models, and develop benchmarks measuring grounded reasoning quality, provenance completeness, and efficiency. Experimental evaluation will compare vector-based RAG, hybrid structured retrieval, and fully KG-driven agentic reasoning across **scientific research discovery and engineering decision-support** domains.

Expected outcomes include the PA-AKG architecture and open research prototype, methods for integrating structured knowledge graphs with generative retrieval pipelines, evaluation benchmarks for grounded reasoning and provenance validation in agentic systems, and reproducible experimental baselines. This work aims to significantly improve the **reliability, transparency, and robustness of agentic AI systems operating over complex real-world knowledge environments**.

References

Aidan Hogan, Eva Blomqvist, Michael Cochez, Claudia d'Amato, Gerard de Melo, Claudio Gutierrez, Sabrina Kirrane, José Emilio Labra Gayo, Roberto Navigli, Sebastian Neumaier, Axel-Cyrille Ngonga Ngomo, Axel Polleres, Sabbir M. Rashid, Anisa Rula, Lukas Schmelzeisen, Juan F. Sequeda, Steffen Staab, and Antoine Zimmermann. Knowledge graphs. *ACM Computing Surveys*, 54(4):1–37, 2021. doi: 10.1145/3447772.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12): 1–38, 2023. doi: 10.1145/3571730.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020.

Joon Sung Park, Joseph C. O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023. doi: 10.1145/3586183.3606763.