

Agentic Knowledge Graphs for AI-Accelerated Engineering with Simulation-Grounded Validation

Kereshmeh Afsari (PI)

ARCADE Lab

Myers-Lawson School of Construction
Virginia Tech

Chris Brown (Co-PI)

Code World No Blanket

Department of Computer Science
Virginia Tech

Jason Cusati (Student Researcher)

Code World No Blanket

Department of Computer Science
Virginia Tech

Primary Topic: AI Accelerated Science & Engineering
Secondary: Agentic AI

I. PROBLEM STATEMENT AND MOTIVATION

AI-driven scientific breakthroughs increasingly rely on tightly coupling learned models with structured knowledge and simulation [1], [2]. Engineering disciplines face an analogous opportunity, yet large language models (LLMs) remain fundamentally misaligned with practice: they hallucinate physical quantities [3], fail to enforce constraints, and lack grounding in the structured ontologies and simulation feedback loops that govern real-world systems. While knowledge-grounded generation—combining retrieval-augmented generation (RAG) [4], knowledge graphs (KGs) [5], and multi-source integration—improves factual consistency, it does not address the central application challenge: no established methodology validates AI-generated engineering outputs against physical laws and regulatory constraints [6]. This trust gap is not academic. In safety-critical engineering, an unverified design can violate building codes such as the American Wood Council’s National Design Specification (AWC-NDS) or the International Building Code (IBC), miss capacity thresholds defined by level-of-service (LOS) limits in the Highway Capacity Manual 2022 (HCM2022), or breach the North American Electric Reliability Corporation Transmission Planning standards (NERC-TPL)—outcomes for which firms carry legal liability and that today’s agents cannot detect, much less prevent. Manufacturer-cycle verification fills this gap by routing every design through weeks-to-months of human-and-simulation review, which keeps construction, transportation, and energy workflows from absorbing agentic AI at the speed Amazon’s downstream services already demand. Amazon’s stake is direct: Bedrock Knowledge Bases [7] and the construction-logistics workflows around AWS-scale supply chains both need agentic outputs that arrive certified, with calibrated uncertainty bounds, on minutes-to-hours latency. Closing this gap requires not better retrieval but a principled, simulation-backed validation methodology that an agent can call in-the-loop.




Fig. 1. **AKG-E Architecture.** Domain modules (Construction: primary; Transportation, Energy: extensions) plug into shared ontology, simulation, and VVUQ layers. The dashed feedback arrow represents iterative constraint-driven refinement.

II. PROPOSED APPROACH: THE AKG-E FRAMEWORK

This project proposes **Agentic Knowledge Graph Engineering (AKG-E)**, a domain-agnostic framework integrating knowledge-grounded reasoning with simulation-in-the-loop Verification, Validation, and Uncertainty Quantification (VVUQ). AKG-E applies three shared layers uniformly across engineering domains—with only the domain knowledge graph and simulator swapped: (1) a *domain ontology layer* encoding engineering KGs with entities, constraints, and physical laws; (2) a *simulation integration layer* exposing domain tools (Finite Element Analysis (FEA) solvers [8], digital twin microsimulators [9], and Partial Differential Equation (PDE) solvers) as typed agent-callable interfaces; and (3) a *VVUQ validation layer* that verifies outputs against governing equations, quantifies uncertainty via Monte Carlo (MC) propagation through agent pipelines, and enforces regulatory constraints. Agents orchestrate across retrieval-augmented generation, structured graph queries, and simulation feedback, extending industry knowledge-grounding approaches (e.g., Amazon-scale KG systems [10]) with physics-based validation.

The contribution that AKG-E adds beyond prior KG + agent systems is the **closed-loop coupling of agent generation with simulation-grounded validation**: candidate outputs are routed through the VVUQ layer before they are returned, and the

layer’s signals—constraint-satisfaction status, identification of the binding limits when constraints fail, and uncertainty bounds on quantitative outputs—feed back into the agent’s next iteration. The specific validation operations are *domain-instantiated* (e.g., physics-based residuals and codified-standard lookups for structural pilots; microsimulation outcomes and level-of-service thresholds for transportation), and the framework’s contribution is the typed, iterable *interface* that couples KG, simulator, and agent—not a fixed catalog of checks. Unlike text-based critic loops or single-step retrieval-validity gates, which can flag inconsistency in prose but cannot reason about physical or regulatory satisfaction, AKG-E’s loop closes against quantitative simulation outputs held to codified engineering standards.

III. RESEARCH PLAN

A. Research Objectives

This project pursues three objectives. **Objective 1 (Framework):** generalize the existing Laminated Veneer Lumber (LVL) VVUQ pipeline into the AKG-E framework—a typed, agent-callable propose $\rightarrow$ VVUQ $\rightarrow$ feedback loop coupling a domain knowledge graph, a simulation tool surface, and a governing-equation / code-compliance / constraint-satisfiability validation gate. **Objective 2 (Validation):** demonstrate AKG-E on two engineering domains (construction: Building Information Modeling / Industry Foundation Classes (BIM/IFC) inputs + FEA against AWC-NDS; transportation: digital-twin microsimulation against HCM2022) and measure the three-tier metrics—correctness, calibration, efficiency—against B0/B1/B2 baselines. **Objective 3 (Dissemination):** release the AKG-E framework, KG ontology templates, the simulator-tool interface specification, and the cross-domain benchmark as open source; publish peer-reviewed results in AI-for-engineering and agentic-AI venues.

B. Methodology

Agent $\leftrightarrow$ VVUQ-gate protocol. The propose-VVUQ-feedback loop is realized as a typed call/response between the agent and the validation layer. The agent emits an artifact specification serialized to the domain ontology schema; the VVUQ layer evaluates the artifact against the physical relations and codified constraints relevant to the domain, then returns structured validation signals—constraint-satisfaction outcomes, the binding limits when constraints fail, and propagated uncertainty bounds on quantitative outputs. The specific evaluations are domain-instantiated: for the construction pilot the layer draws on the existing LVL VVUQ work (FEniCS [8]-based structural relations with codified-standard lookups), and the transportation pilot composes microsimulation outcomes with level-of-service thresholds. The framework’s contribution is the *interface specification*—typed inputs, structured signals, propagated uncertainty—not a fixed catalog of which checks must run; domains compose their own check set against the shared signal contract.

Uncertainty propagation. MC sampling is applied jointly over (a) simulator inputs—tributary widths, distributed loads,

modulus of elasticity, demand profiles, grid contingency states—modeled as random variables with priors elicited from codified design tables and prior-art reliability literature, and (b) agent sampling temperature, so that the propagated 95% confidence interval (CI) captures both the simulator’s epistemic uncertainty and the LLM’s stochasticity. For each artifact, $N = 1000$ MC samples is the default; convergence is tested against $N = 10,000$ during verification, with all draws traced through the KG to preserve provenance.

Construction-pilot specifics. The pilot reuses three already-implemented modules from the construction-ai prototype: a BIM/IFC parser [11] that lifts wall, opening, and load-path entities into the KG; a typed FEA tool-call surface that wraps FEniCS solvers behind a stable agent-callable interface; and a refinement step that closes the loop by translating gate-emitted violations (e.g., “deflection 0.42 in exceeds $L/240 = 0.40$ in”) back into KG-conditioned prompts for the next iteration. The LVL VVUQ pipeline already operational in construction-ai supplies the methodology anchor and the closed-form verification baseline; AKG-E generalizes its Euler–Bernoulli + Monte Carlo pattern into the iterable propose-validate-refine protocol above.

C. Evaluation Plan

The evaluation enforces a three-tier metric framework that refines the abstract’s constraint-violation reduction target into a more rigorous structure: correctness (the abstract’s $>50\%$ violation-reduction target preserved as the must-pass floor), calibration, and efficiency. **Correctness** is the non-negotiable floor: *constraint-violation rate* measured against codified standards—AWC-NDS / IBC2021 (construction), HCM2022 LOS thresholds (transportation), NERC-TPL contingency criteria (energy, designed extension)—must approach zero on the benchmark and be statistically lower than every baseline. Agents that are fast and wrong are unfundable in safety-critical engineering. **Calibration:** the empirical coverage probability of the propagated 95% CI must contain the ground truth in $\geq 95\%$ of benchmark tasks; calibration is validated against FEA reference designs (construction domain, via the LVL VVUQ pipeline outputs) and microsimulation reference outcomes (transportation domain, via the existing Virginia Tech Transportation Institute (VTTI) digital-twin proposal). **Efficiency** (headline contribution): wall-clock time-to-validated-decision per task is compared to the status-quo manufacturer-cycle baseline (weeks-to-months for construction VVUQ done by manufacturers); AKG-E targets minutes-to-hours per task at codified-standard correctness with calibrated UQ.

Baselines (three-tier ablation): **B0** Claude alone (LLM-only; no retrieval, no validation loop); **B1** Claude + graph-RAG over the curated domain KG (retrieval without the agentic VVUQ loop); **B2** Claude + graph-RAG + agentic VVUQ loop (full AKG-E—this is the contribution). Each tier is run on the same benchmark, identical retrieval index, and identical prompt scaffold up to the loop boundary, so the marginal effect of the VVUQ loop is isolated.

Benchmark suite. The funded benchmark covers two domains, with a third designed in for future extension. *Construction (funded)*: header sizing and material takeoff with AWC-NDS code-compliance checks; ground truth from the LVL VVUQ pipeline and AWC-NDS reference designs. *Transportation (funded)*: signal-timing and digital-twin scenarios drawn from the existing VTTI Virginia Transportation Safety Index (VTSI) proposal; ground truth from microsimulation outcomes and HCM LOS targets. *Energy (designed extension)*: the three AKG-E layers apply to grid contingency analysis (NERC-TPL criteria) by swapping only the domain KG and PDE-solver interface; a proof-of-concept (POC) task is scheduled for Months 4–12 to demonstrate the domain-swap claim empirically. **Statistical analysis:** $n = 5$ seeds minimum per (domain $\times$ baseline) cell; 95% bootstrap confidence intervals on all three metrics; significance threshold $p < 0.05$ for headline claims. All three metrics are reported for every cell—no cherry-picking which metric to feature per condition.

IV. TIMELINE, DELIVERABLES, AND RISKS

The funded year begins **August 2026** and is structured as a 3-month aggressive build (M1–M3, Aug–Oct 2026) followed by 9 months of dissemination, robustness, scaling, and the energy-domain extension (M4–M12, Nov 2026–Jul 2027). Table I maps the workstreams onto the calendar.

Build phase rationale (M1–M3). The construction-pilot anchor reuses the existing LVL VVUQ pipeline (BIM/IFC parser, FEniCS tool-call surface, Monte Carlo refinement), so M1 *generalizes* existing code rather than green-fielding the framework; the transportation pilot lands per the existing VTTI digital-twin proposal across M2–M4; the cross-domain benchmark begins on construction and adds transportation as the second pilot stabilizes, with statistical analysis and the first paper draft completing in M5–M6. *Extension phase (M4–M12).* The energy domain-swap POC (NERC-TPL KG + PDE-solver interface) empirically validates the framework’s extensibility claim across M5–M8 by swapping only the domain KG and simulator without modifying the shared ontology, simulation-integration, or VVUQ-validation layers. Robustness experiments stress-test model swaps and prompt perturbations; AWS Bedrock scaling studies establish the deployment surface; the final quarter delivers publications and the open-source release. *Deliverables:* (1) the open-source AKG-E framework; (2) domain-adaptable KG ontology templates and the typed FEA / digital-twin / PDE simulator-tool interface specification; (3) the cross-domain benchmark with ground-truth datasets (construction + transportation funded, energy POC for the extensibility demonstration); (4) peer-reviewed publications targeting AI-for-engineering and agentic-AI venues.

Risks and Mitigation. (a) *Agentic VVUQ loop slow to converge.* The propose-validate-refine loop may iterate beyond budget on edge cases. Mitigation: cap iterations at $K = 10$ per artifact and accept best-so-far with explicit UQ caveats; this preserves correctness and calibration metrics while exposing efficiency as the variable under study. (b) *Code-compliance ontology completeness.* The AWC-NDS and HCM2022 corpora

are large; if the curated KG ontology misses provisions, the gate may undercount violations. Mitigation: prioritize the highest-frequency provisions in the benchmark suite (header sizing for construction; uninterrupted-flow segments for transportation) and document ontology-coverage scope explicitly so reviewers can interpret reported violation rates against a stated provision boundary; coverage expands in Months 4–6 before publication. (c) *Benchmark ground-truth coverage thinner than expected.* If curating AWC-NDS reference designs or HCM-validated microsimulation scenarios takes longer than budgeted, mitigation is to ship a minimum-viable product (MVP) benchmark exercising B0/B1/B2 across a reduced task taxonomy in Month 3 and expand the suite in Months 4–6 before publication; the three-tier metrics remain reportable on the MVP, and the energy extension’s KG-schema design proceeds in parallel rather than after.

V. TEAM AND CAPABILITIES

Dr. Kereshmeh Afsari (PI) directs the ARCADE Lab at the Myers-Lawson School of Construction, with deep expertise in BIM, 3D visualization, and the construction-AI initiative that supplies AKG-E’s primary domain testbed; she owns the ARCADE infrastructure (BIM/IFC datasets, simulation pipelines, AWS account) and the construction-industry partnerships that anchor the LVL pilot. **Dr. Chris Brown (Co-PI)** directs the Code World No Blanket lab in the Department of Computer Science at Virginia Tech and contributes expertise in software engineering tooling and the empirical evaluation of AI-generated code, including the Brown–Cusati line of work on the SE beliefs of generative AI tools [12] and on knowledge accumulation in software engineering [13]—both pointing to the verification and grounding weaknesses AKG-E sets out to close; he advises the student-research thread on agentic knowledge graphs. **Jason Cusati (Student Researcher)** brings the agentic-KG research-progression substrate that supplies AKG-E’s orchestration layer, the VTTI Virginia Transportation Safety Index calculator (and its digital-twin extension proposal) that anchors the transportation domain, and direct VVUQ training under Dr. Christopher Roy’s graduate coursework (Oberkampf–Roy formulation [6]).

Two Graduate Research Assistants full-time during the 3-month build provide the engineering bandwidth. Three working prototypes ground the feasibility argument and align directly with the build plan: (1) the LVL VVUQ pipeline (Euler–Bernoulli + Monte Carlo over BIM/IFC inputs) already operational in the construction-ai system, supplying the methodology anchor for Month 1’s generalization step; (2) the construction-ai KG + multi-agent material-takeoff web application, supplying the agent-orchestration substrate; and (3) the existing VTTI transportation digital-twin proposal, supplying the Month 2 implementation target. ARCADE Lab provides BIM/IFC reference datasets, FEA simulation pipelines, and the AWS account that hosts the benchmark infrastructure.

TABLE I

PROJECT TIMELINE. **M1–M3** IS THE AGGRESSIVE BUILD PHASE; **M4–M12** DELIVERS DISSEMINATION, ROBUSTNESS, SCALING, AND THE ENERGY DOMAIN-SWAP PROOF-OF-CONCEPT THAT DEMONSTRATES AKG-E’S EXTENSIBILITY CLAIM.

Activity	Month (Aug 2026 – Jul 2027)											
	M1	M2	M3	M4	M5	M6	M7	M8	M9	M10	M11	M12
Framework: agentic propose→VVUQ→feedback loop	■	■	■									
Construction pilot: LVL VVUQ generalization (BIM/IFC + FEA)	■	■	■									
Transportation pilot: digital-twin microsim integration		■	■	■								
Benchmark suite: B0 / B1 / B2 × 2 domains × $n \geq 5$ seeds			■	■	■							
Statistical analysis + first paper draft					■	■						
Energy extension: NERC-TPL KG + PDE-solver interface (POC)					■	■	■	■				
Robustness experiments (model swaps, prompt perturbations)							■	■	■	■		
Scaling studies on AWS Bedrock								■	■	■	■	■
Publications + open-source release										■	■	■

VI. INTELLECTUAL MERIT AND AMAZON FIT

Intellectual merit: AKG-E contributes a principled VVUQ methodology for agentic AI outputs—a generalizable result that separates domain knowledge from shared orchestration and validation, and that is testable, falsifiable, and reusable across engineering disciplines. **Broader impact:** certified AI for safety-critical engineering—reducing hallucinated designs, enforcing code compliance, and enabling principled human-in-the-loop thresholds in workflows where errors carry legal and physical consequences. **Amazon fit:** AKG-E targets deployment via Amazon Bedrock Knowledge Bases [7]; the construction-logistics workflows around AWS-scale supply chains are the most direct downstream consumer; and the AWS-scale data pipelines the framework integrates with make AKG-E immediately relevant to Amazon’s applied AI and construction-logistics roadmap.

REFERENCES

- [1] J. Abramson, J. Adler, J. Dunger, R. Evans, T. Green, A. Pritzel, O. Ronneberger, L. Willmore, A. J. Ballard, J. Bambrick *et al.*, “Accurate structure prediction of biomolecular interactions with AlphaFold 3,” *Nature*, vol. 630, pp. 493–500, 2024.
- [2] A. Merchant, S. Batzner, S. S. Schoenholz, M. Aykol, G. Cheon, and E. D. Cubuk, “Scaling deep learning for materials discovery,” *Nature*, vol. 624, pp. 80–85, 2023.
- [3] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. Bang, A. Madotto, and P. Fung, “Survey of hallucination in natural language generation,” *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, 2023.
- [4] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 9459–9474.
- [5] D. Edge, H. Trinh, N. Cheng, J. Bradley, A. Chao, A. Mody, S. Truitt, and J. Larson, “From local to global: A graph RAG approach to query-focused summarization,” arXiv preprint arXiv:2404.16130, 2024. [Online]. Available: <https://arxiv.org/abs/2404.16130>
- [6] W. L. Oberkampff and C. J. Roy, *Verification and Validation in Scientific Computing*. Cambridge University Press, 2010.
- [7] Amazon Web Services, “Amazon Bedrock knowledge bases now supports hybrid search,” AWS Machine Learning Blog, Mar. 2024, accessed 2026-03-10. [Online]. Available: <https://aws.amazon.com/blogs/machine-learning/amazon-bedrock-knowledge-bases-now-supports-hybrid-search/>
- [8] A. Logg, K.-A. Mardal, and G. N. Wells, Eds., *Automated Solution of Differential Equations by the Finite Element Method: The FEniCS Book*. Springer, 2012.
- [9] F. Tao, H. Zhang, A. Liu, and A. Y. C. Nee, “Digital twin in industry: State-of-the-art,” *IEEE Transactions on Industrial Informatics*, vol. 15, no. 4, pp. 2405–2415, 2019.
- [10] C. Yu, X. Liu, J. Maia, T. Cao, Y. Li, Y. Gao, Y. Song, R. Goutam, H. Zhang, B. Yin, and Z. Li, “COSMO: A large-scale e-commerce common sense knowledge generation and serving system at amazon,” in *Companion Proceedings of the ACM International Conference on Management of Data (SIGMOD-Companion)*. ACM, 2024, pp. 386–399.
- [11] C. Eastman, P. Teicholz, R. Sacks, and K. Liston, *BIM Handbook: A Guide to Building Information Modeling for Owners, Managers, Designers, Engineers and Contractors*, 2nd ed. John Wiley & Sons, 2011.
- [12] C. Brown and J. Cusati, “Exploring the evidence-based se beliefs of generative ai tools,” in *2025 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM)*. IEEE, 2025, pp. 386–392.
- [13] J. Cusati and C. Brown, “From papers to progress: Rethinking knowledge accumulation in software engineering,” in *To Appear: 2026 IEEE/ACM 48th International Conference on Software Engineering: Future of Software Engineering (ICSE-FoSE)*, 2026.